

MACHINE LEARNING (IMC-4302C)

LAB 1: LINEAR REGRESSION

School Year
2017-2018

6/2/2018

About your tutor and labs

- Slim BEN AMOR: Slim.ben-amor@inria.fr
 - Bachelor from Ecole Polytechnique of Tunisia + Master from Paris-Saclay
 - Phd Student at INRIA
 - Online courses achievements (Machine Learning from Stanford and AI from Berkley University)
- 10 Labs:
 - 2 labs per week
 - 1 report per week (your code for 2 labs + 2-3 pages for your comments/observations on each question)
 - Grade: 5 reports 50% + Final project 50%

Notation

$$X = \begin{bmatrix} x_{1,0} & x_{1,1} & \dots & x_{1,n-1} \\ x_{2,0} & x_{2,1} & \dots & x_{2,n-1} \\ \vdots & \vdots & \dots & \vdots \\ x_{m,0} & x_{m,1} & \dots & x_{m,n-1} \end{bmatrix}$$

Where: - m is number of samples
 - n is number of features including
 y-intercept (bias)

$$x_i = [x_{i,0} \quad x_{i,1} \quad \dots \quad x_{i,n-1}]$$

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_{n-1} \end{bmatrix}$$

$$h_{\theta}(x_i) = \theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2} + \dots + \theta_{n-1} x_{i,n-1}$$

$$h_{\theta}(X) = X\theta$$

Cost Function: Mean Square Error (MSE)

- Cost Function
$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i)^2$$
- Partial derivative
$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i) x_j$$
$$j = 0 \dots n - 1$$

Minimizing Cost Function: Gradient descent Algorithm

- Update equation for parameters θ_j

Where : α step or learning rate

$$\theta_j = \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}$$
$$j = 0 \dots n - 1$$

- The update of θ_j should be made simultaneously. Because if we update θ_1 the vector θ will change. Hence, when updating θ_2 we will calculate the partial derivative of $J(\theta)$ on a different point θ as the initial θ .

- Simultaneously update

$$\theta_j^{temp} = \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j} \text{ for all } j = 0 \dots n - 1$$
$$\text{then } \theta_j = \theta_j^{temp}$$

Vectorized Implementation

- Cost function

$$J(\theta) = \frac{1}{2m}(X\theta - y)^\top (X\theta - y)$$

- Cost function gradient

$$\nabla J(\theta) = \frac{1}{m}X^\top (X\theta - y)$$

- Gradient descent update

$$\theta = \theta - \alpha \nabla J(\theta)$$

Feature normalization

- Cost function
$$x_{norm} = \frac{(x - \bar{x})}{\sigma_x}$$

Where : \bar{x} is the mean of x
 σ_x is the standard deviation